


2) Logistic Regression

Useful for binary classification. Should NOT be used for regression



If we will use simple linear regression, it can misclassify data (cause of points/outliers far away) like so:



So we need to weight the samples (that are far away) much lower. That's where the logistic function (ie sigmoid function) comes into play

Sigmoid function

It looks as follows:



$$y_{\beta}(x) = \frac{1}{1+e^{-(\beta_0 + \beta_1 x + \varepsilon)}}$$

So rather than the slope being $mx+b$, it is $1 / (1 + e^{-(mx+b)})$

$1 / (1 + e^{-(mx+b)})$ can be further represented as

$$P(x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x + \varepsilon)}} \quad \Rightarrow \quad P(x) = \frac{e^{(\beta_0 + \beta_1 x)}}{1 + e^{(\beta_0 + \beta_1 x)}}$$

Logistic Function

$$\frac{P(x)}{1 - P(x)} = e^{(\beta_0 + \beta_1 x)}$$

Odds Ratio

and the odds ratio can further be represented as

$$\log \left[\frac{P(x)}{1 - P(x)} \right] = \beta_0 + \beta_1 x$$


Odds Ratio

One vs all

This is a method used when we are not dealing with binary classes. Here, for every class, we put that class's points in a separate class and all the other points in another class


So for example, suppose we have 3 classes (a,b,c). So first, we will train a model with all points of **a** being put in **class1** and points of **b** and **c** being put in **class2**. This model will represent the probability that a given point will be of the **a** class. Similarly, we will train two more models, one representing the probability of a point being of class **b** and the other representing the probability of a point being of class **c**. We will then make our final prediction based on which model gave us the highest probability

For example, in a dataset like this:




We end up with the following models:


(blue vs purple+pink)



(purple vs blue+pink)



(pink vs blue+purple)



Code

In code, we can pass hyper parameters like `penalty`, `c`, and `solver`

`penalty` is whether we want to use l1, l2, or elastic net regularisation

`c` is the inverse of λ

There is also the `LogisticRegressionCV` class which tunes regularisation parameters via cross validation